

```
hbase(main):001:0> create 'flight','finfo','fsch'  
0 row(s) in 1.0600 seconds
```

```
=> Hbase::Table - flight  
hbase(main):002:0> list  
TABLE  
employee  
flight  
2 row(s) in 0.1130 seconds
```

```
=> ["employee", "flight"]  
hbase(main):003:0> put 'flight','1','finfo:source','Pune'  
0 row(s) in 0.3450 seconds
```

```
hbase(main):004:0> put 'flight','1','finfo:destination','Mumbai'  
0 row(s) in 0.0170 seconds
```

```
hbase(main):005:0> put 'flight','1','fsch:depar','10.00a.m'  
0 row(s) in 0.0270 seconds
```

```
hbase(main):006:0> put 'flight','1','fsch:arrival','11.00a.m'  
0 row(s) in 0.0260 seconds
```

```
hbase(main):008:0> put 'flight','2','finfo:source','Pune'  
0 row(s) in 0.0170 seconds
```

```
hbase(main):009:0> put 'flight','2','finfo:destination','Kolkata'  
0 row(s) in 0.0170 seconds
```

```
hbase(main):013:0> put 'flightdb','2','fsh:depar','1:00 am'  
0 row(s) in 0.0080 seconds
```

```
hbase(main):014:0> put 'flightdb','3','fin:source','Kolkata'  
0 row(s) in 0.0120 seconds
```

```
hbase(main):015:0> put 'flightdb','3','fin:destination','Goa'  
0 row(s) in 0.0190 seconds
```

```
hbase(main):016:0> put 'flightdb','3','fsh:arrival','10:00 pm'  
0 row(s) in 0.0100 seconds
```

```
hbase(main):017:0> put 'flightdb','3','fsh:depar','8:00 pm'  
0 row(s) in 0.0090 seconds
```

```
hbase(main):018:0> put 'flightdb','4','fin:source','Mumbai'  
0 row(s) in 0.0310 seconds
```

```
hbase(main):019:0> put 'flightdb','4','fin:destination','Chennai'  
0 row(s) in 0.0080 seconds
```

```
hbase(main):020:0> put 'flightdb','4','fsh:arrival','7:00 pm'
```

0 row(s) in 0.0080 seconds

hbase(main):021:0> put 'flightdb','4','fsh:depar','6:00 pm'
0 row(s) in 0.0420 seconds

hbase(main):022:0> put 'flightdb','5','fin:source','Mysore'
0 row(s) in 0.0080 seconds

hbase(main):023:0> put 'flightdb','5','fin:destination','ooty'
0 row(s) in 0.0060 seconds
hbase(main):024:0> put 'flightdb','5','fsh:arrival','1:00 pm'
0 row(s) in 0.0090 seconds

hbase(main):025:0> put 'flightdb','5','fsh:depar','12:30 pm'
0 row(s) in 0.0120 seconds

hbase(main):026:0> put 'flightdb','6','fin:source','Chennai'
0 row(s) in 0.0110 seconds

hbase(main):027:0> put 'flightdb','6','fin:destination','Gujrat'
0 row(s) in 0.0070 seconds
hbase(main):006:0> put 'flightdb','6','fsh:arrival','1.00 pm'
0 row(s) in 0.0080 seconds

hbase(main):007:0> put 'flightdb','6','fsh:depar','12.30 pm'
0 row(s) in 0.0090 seconds

hbase(main):008:0> put 'flightdb','7','fin:source','3.30 pm'
0 row(s) in 0.0100 seconds

hbase(main):009:0> put 'flightdb','7','fin:source','Punjab'
0 row(s) in 0.0300 seconds

hbase(main):010:0> put 'flightdb','7','fin:destination','Rajasthan'
0 row(s) in 0.0090 seconds
hbase(main):011:0> put 'flightdb','7','fsh:arrival','9.00 pm'
0 row(s) in 0.0170 seconds

hbase(main):012:0> put 'flightdb','7','fsh:depar','7:00 pm'
0 row(s) in 0.0080 seconds

hbase(main):013:0> put 'flightdb','8','fin:source','Haryana'
0 row(s) in 0.0190 seconds

hbase(main):014:0> put 'flightdb','8','fin:destination','Chandigarh'
0 row(s) in 0.0070 seconds

hbase(main):015:0> put 'flightdb','8','fsh:depar','10:00 pm'
0 row(s) in 0.0090 seconds

hbase(main):016:0> put 'flightdb','8','fsh:depar','9:00 pm'
0 row(s) in 0.0250 seconds

```
hbase(main):017:0> put 'flightdb','9','fin:source','MP'
```

```
0 row(s) in 0.0100 seconds
```

```
hbase(main):018:0> put 'flightdb','9','fin:destination','UP'
```

```
0 row(s) in 0.0130 seconds
```

```
hbase(main):019:0> put 'flightdb','8','fsh:arrival','10:00 pm'
```

```
0 row(s) in 0.0360 seconds
```

```
hbase(main):020:0> put 'flightdb','9','fsh:arrival','10:00 am'
```

```
0 row(s) in 0.0180 seconds
```

```
hbase(main):021:0> put 'flightdb','9','fsh:depar','9:00 am'
```

```
0 row(s) in 0.0090 seconds
```

```
hbase(main):022:0> put 'flightdb','10','fin:source','Patna'
```

```
0 row(s) in 0.0080 seconds
```

```
hbase(main):023:0> put 'flightdb','10','fin:destination','Ranchi'
```

```
0 row(s) in 0.0060 seconds
```

```
hbase(main):024:0> put 'flightdb','10','fsh:arrival','11:00 am'
```

```
0 row(s) in 0.0080 seconds
```

```
hbase(main):025:0> put 'flightdb','10','fsh:depar','10:00 am'
```

```
0 row(s) in 0.0070 seconds
```

```
hbase(main):099:0> exit
```

```
[cloudera@quickstart ~]$ hive
```

```
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
```

```
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
```

```
hive> beeline -u jdbc:hive2://localhost:10000
```

```
> CREATE EXTERNAL TABLE flight_hive (
```

```
>   rowkey STRING,
```

```
>   source STRING,
```

```
>   destination STRING,
```

```
>   arrival STRING,
```

```
>   departure STRING
```

```
> )
```

```
> STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
```

```
> WITH SERDEPROPERTIES (
```

```
>   "hbase.columns.mapping" = ":key,fin:source,fin:destination,fsh:arrival,fsh:departure"
```

```
> )
```

```
> TBLPROPERTIES ('hbase.table.name'='flightdb');
```

```
OK
```

```
Time taken: 4.52 seconds
```

```
hive> SHOW TABLES;
```

```
OK
```

```
employee
```

```
flight_hive
```

```
Time taken: 0.563 seconds, Fetched: 2 row(s)
```

```

hive> DESC flight_hive;
OK
rowkey                string          from deserializer
source                string          from deserializer
destination            string          from deserializer
arrival               string          from deserializer
departure              string          from deserializer
Time taken: 0.391 seconds, Fetched: 5 row(s)
hive> SELECT
  >  AVG(UNIX_TIMESTAMP(arrival, 'hh:mm a')) AS avg_arrival_seconds,
  >  AVG(UNIX_TIMESTAMP(departure, 'hh:mm a')) AS avg_departure_seconds
  > FROM flight_hive;
Query ID = cloudera_20250326000808_e5b290a3-48fa-4b7d-a1e9-910555f7153b
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1742968947688_0001, Tracking URL =
http://quickstart.cloudera:8088/proxy/application_1742968947688_0001/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1742968947688_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2025-03-26 00:09:23,718 Stage-1 map = 0%, reduce = 0%
2025-03-26 00:09:39,678 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.91 sec
2025-03-26 00:09:55,986 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.24 sec
MapReduce Total cumulative CPU time: 3 seconds 240 msec
Ended Job = job_1742968947688_0001
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 3.24 sec HDFS Read: 15722 HDFS Write:
11 SUCCESS
Total MapReduce CPU Time Spent: 3 seconds 240 msec
OK
78300.0      NULL
Time taken: 58.334 seconds, Fetched: 1 row(s)
hive> SELECT
  >  FROM_UNIXTIME(AVG(UNIX_TIMESTAMP(arrival, 'hh:mm a')), 'hh:mm a') AS
avg_arrival_time,
  >  FROM_UNIXTIME(AVG(UNIX_TIMESTAMP(departure, 'hh:mm a')), 'hh:mm a') AS
avg_departure_time
  > FROM flight_hive;
hive> SELECT arrival, departure FROM flight_hive LIMIT 10;
OK
11:00 am     NULL
11:00 am     NULL
2:00 am      NULL
10:00 pm     NULL
7:00 pm      NULL
1:00 pm      NULL

```

1.00 pm NULL
9.00 pm NULL
10:00 pm NULL
10:00 am NULL

Time taken: 2.553 seconds, Fetched: 10 row(s)

hive> SELECT

> FROM_UNIXTIME(AVG(UNIX_TIMESTAMP(arrival, 'hh:mm a')), 'hh:mm a') AS

avg_arrival_time,

> FROM_UNIXTIME(AVG(UNIX_TIMESTAMP(departure, 'hh:mm a')), 'hh:mm a') AS

avg_departure_time

> FROM flight_hive

> WHERE arrival IS NOT NULL AND departure IS NOT NULL;

SELECT

FROM_UNIXTIME(AVG(UNIX_TIMESTAMP(arrival, 'hh:mm a')), 'hh:mm a') AS

avg_arrival_time,

FROM_UNIXTIME(AVG(UNIX_TIMESTAMP(departure, 'hh:mm a')), 'hh:mm a') AS

avg_departure_time

FROM flight_hive;

avg_arrival_time avg_departure_time

3:06 pm 2.05 p