

DSBDA MOCK - Hive Inner Join on Employees and Department
Roll no - 33322

[cloudera@quickstart Desktop]\$ hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
hive> create table employee(emp_id int, emp_name string, dept_id int) row format delimited fields terminated by ','

> ;

OK

Time taken: 5.686 seconds

hive> create table employees(emp_id int, emp_name string, dept_id int) row format delimited fields terminated by ',' stored as TEXTFILE;

OK

Time taken: 0.323 seconds

hive> create Display all 461 possibilities? (y or n)^C[cloudera@quickstart Desktop]\$ ^C

[cloudera@quickstart Desktop]\$ hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
hive> create table departments(dept_name string, dept_id int) row format delimited fields terminated by ',';

OK

Time taken: 1.991 seconds

hive> create table departments(dept_name string, dept_id int) row format delimited fields terminated by ',' stored as TEXTFILE;

FAILED: Execution Error, return code 1 from org.apache.hadoop.hive.ql.exec.DDLTask.
AlreadyExistsException(message:Table departments already exists)

hive> DROP TABLE departments;

OK

Time taken: 1.259 seconds

hive>

> create table departments(dept_name string, dept_id int) row format delimited fields terminated by ',' stored as TEXTFILE;

OK

Time taken: 0.144 seconds

hive> load data inpath 'HIVEJOIN/employees.txt' into table employees;

Loading data to table default.employees

chgrp: changing ownership of

'hdfs://quickstart.cloudera:8020/user/hive/warehouse/employees/employees.txt': User does not belong to hive

Table default.employees stats: [numFiles=1, totalSize=43]

OK

Time taken: 1.483 seconds

hive> load data inpath 'HIVEJOIN/departments.txt' into table departments;

Loading data to table default.departments

```

chgrp: changing ownership of
'hdfs://quickstart.cloudera:8020/user/hive/warehouse/departments/departments.txt': User
does not belong to hive
Table default.departments stats: [numFiles=1, totalSize=15]
OK
Time taken: 0.555 seconds
hive> hive> load data inpath 'HIVEJOIN/departments.txt' into table departments;
NoViableAltException(26@[])
    at org.apache.hadoop.hive.ql.parse.HiveParser.statement(HiveParser.java:1020)
    at org.apache.hadoop.hive.ql.parse.ParseDriver.parse(ParseDriver.java:199)
    at org.apache.hadoop.hive.ql.parse.ParseDriver.parse(ParseDriver.java:166)
    at org.apache.hadoop.hive.ql.Driver.compile(Driver.java:393)
    at org.apache.hadoop.hive.ql.Driver.compile(Driver.java:307)
    at org.apache.hadoop.hive.ql.Driver.compileInternal(Driver.java:1110)
    at org.apache.hadoop.hive.ql.Driver.runInternal(Driver.java:1158)
    at org.apache.hadoop.hive.ql.Driver.run(Driver.java:1047)
    at org.apache.hadoop.hive.ql.Driver.run(Driver.java:1037)
    at org.apache.hadoop.hive.cli.CliDriver.processLocalCmd(CliDriver.java:207)
    at org.apache.hadoop.hive.cli.CliDriver.processCmd(CliDriver.java:159)
    at org.apache.hadoop.hive.cli.CliDriver.processLine(CliDriver.java:370)
    at org.apache.hadoop.hive.cli.CliDriver.executeDriver(CliDriver.java:756)
    at org.apache.hadoop.hive.cli.CliDriver.run(CliDriver.java:675)
    at org.apache.hadoop.hive.cli.CliDriver.main(CliDriver.java:615)
    at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
    at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:57)
    at
sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
    at java.lang.reflect.Method.invoke(Method.java:606)
    at org.apache.hadoop.util.RunJar.run(RunJar.java:221)
    at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
FAILED: ParseException line 1:0 cannot recognize input near 'hive' '>' 'load'
hive> Loading data to table default.departments
    > chgrp: changing ownership of
'hdfs://quickstart.cloudera:8020/user/hive/warehouse/departments/departments.txt': User
does not belong to hive
    > Table default.departments stats: [numFiles=1, totalSize=15]
    > OK
    > Time taken: 0.555 seconds
    > hive>
    > [cloudera@quickstart Desktop]$ hive

```

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties

WARNING: Hive CLI is deprecated and migration to Beeline is recommended.

```
hive> select * from employees;
```

OK

```

1      john      2
2      jacob     2
3      tim       1

```

```
4      tom      3
5      liz      1
```

Time taken: 3.118 seconds, Fetched: 5 row(s)

```
hive> select * from departments
```

```
> ;
```

OK

```
hr      1
it      2
qs      3
```

Time taken: 0.151 seconds, Fetched: 3 row(s)

```
hive> select employees.emp_id, employees.emp_name, departments.dept_name from
employees join departments on employees.dept_id = departments.dept_id;
```

Query ID = cloudera_20250325232727_32129e88-1c69-48cc-b40d-97716630d68f

Total jobs = 1

Execution log at:

/tmp/cloudera/cloudera_20250325232727_32129e88-1c69-48cc-b40d-97716630d68f.log

2025-03-25 11:27:48 Starting to launch local task to process map join; maximum memory
= 1013645312

2025-03-25 11:27:52 Dump the side-table for tag: 1 with group count: 3 into file:

file:/tmp/cloudera/2df66564-4086-4dbd-8395-5711558cbcd5/hive_2025-03-25_23-27-33_91
8_9042112746110761413-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile01--.hashtable

2025-03-25 11:27:52 Uploaded 1 File to:

file:/tmp/cloudera/2df66564-4086-4dbd-8395-5711558cbcd5/hive_2025-03-25_23-27-33_91
8_9042112746110761413-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile01--.hashtable
(326 bytes)

2025-03-25 11:27:52 End of local task; Time Taken: 4.325 sec.

Execution completed successfully

MapredLocal task succeeded

Launching Job 1 out of 1

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job_1742969079981_0001, Tracking URL =

http://quickstart.cloudera:8088/proxy/application_1742969079981_0001/

Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1742969079981_0001

Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0

2025-03-25 23:28:23,562 Stage-3 map = 0%, reduce = 0%

2025-03-25 23:28:42,382 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 1.45 sec

MapReduce Total cumulative CPU time: 1 seconds 450 msec

Ended Job = job_1742969079981_0001

MapReduce Jobs Launched:

Stage-Stage-3: Map: 1 Cumulative CPU: 1.45 sec HDFS Read: 6098 HDFS Write: 48

SUCCESS

Total MapReduce CPU Time Spent: 1 seconds 450 msec

OK

```
1      john    it
2      jacob   it
3      tim     hr
4      tom     qs
5      liz     hr
```

Time taken: 69.673 seconds, Fetched: 5 row(s)
hive>

```
[cloudera@quickstart Desktop]$ hdfs dfs -put employees.txt HIVE
[cloudera@quickstart Desktop]$ hdfs dfs -put departments.txt HIVE
put: `HIVE': File exists
[cloudera@quickstart Desktop]$ hdfs dfs -put employees.txt HIVE/
put: `HIVE': File exists
[cloudera@quickstart Desktop]$ hdfs dfs -put employees.txt HIVEJOIN/
put: `HIVEJOIN/': No such file or directory
[cloudera@quickstart Desktop]$ hdfs dfs -mkdir HIVEJOIN
[cloudera@quickstart Desktop]$ hdfs dfs -put employees.txt HIVEJOIN/
[cloudera@quickstart Desktop]$ hdfs dfs -put departments.txt HIVEJOIN/
[cloudera@quickstart Desktop]$
```

Employees.txt

1,john,2
2,jacob,2
3,tim,1
4,tom,3
5,liz,1

Departments.txt

hr,1
it,2
qs,3