Roll No : 33306
DSBDA MOCK:

Problem Statement :
Write an application using HiveQL for Employee information system which will include
a. Creating, Dropping, and altering Database tables.
b. Creating an external Hive table.
c. Load table with data, insert new values and field in the table, Join tables with Hive
d. Create index on Flight Information Table
e. Find the max salary of empoyee in year

```
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
hive> create database mock
    > [cloudera@quickstart ~]$
[cloudera@quickstart ~]$ hiveql
bash: hiveql: command not found
[cloudera@quickstart ~]$ hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j.properties
WARNING: Hive CLI is deprecated and migration to Beeline is recommended.
hive> create databse empdb;
NoViableAltException(26@[682:1: ddlStatement : ( createDatabaseStatement |
switchDatabaseStatement | dropDatabaseStatement | createTableStatement | dropTableStatement |
truncateTableStatement | alterStatement | descStatement | showStatement | metastoreCheck |
createViewStatement | dropViewStatement | createFunctionStatement | createMacroStatement |
createIndexStatement | dropIndexStatement | dropFunctionStatement | dropMacroStatement |
analyzeStatement | lockStatement | unlockStatement | lockDatabase | unlockDatabase |
createRoleStatement | dropRoleStatement | grantPrivileges | revokePrivileges | showGrants |
showRoleGrants | showRolePrincipals | showRoles | grantRole | revokeRole | setRole |
showCurrentRole );])
        at org.antlr.runtime.DFA.noViableAlt(DFA.java:158)
        at org.antlr.runtime.DFA.predict(DFA.java:144)
        at org.apache.hadoop.hive.ql.parse.HiveParser.ddlStatement(HiveParser.java:2301)
        at org.apache.hadoop.hive.ql.parse.HiveParser.execStatement(HiveParser.java:1579)
        at org.apache.hadoop.hive.ql.parse.HiveParser.statement(HiveParser.java:1057)
        at org.apache.hadoop.hive.ql.parse.ParseDriver.parse(ParseDriver.java:199)
        at org.apache.hadoop.hive.ql.parse.ParseDriver.parse(ParseDriver.java:166)
        at org.apache.hadoop.hive.ql.Driver.compile(Driver.java:393)
        at org.apache.hadoop.hive.ql.Driver.compile(Driver.java:307)
        at org.apache.hadoop.hive.ql.Driver.compileInternal(Driver.java:1110)
        at org.apache.hadoop.hive.ql.Driver.runInternal(Driver.java:1158)
        at org.apache.hadoop.hive.ql.Driver.run(Driver.java:1047)
        at org.apache.hadoop.hive.ql.Driver.run(Driver.java:1037)
        at org.apache.hadoop.hive.cli.CliDriver.processLocalCmd(CliDriver.java:207)
        at org.apache.hadoop.hive.cli.CliDriver.processCmd(CliDriver.java:159)
        at org.apache.hadoop.hive.cli.CliDriver.processLine(CliDriver.java:370)
        at org.apache.hadoop.hive.cli.CliDriver.executeDriver(CliDriver.java:756)
        at org.apache.hadoop.hive.cli.CliDriver.run(CliDriver.java:675)
        at org.apache.hadoop.hive.cli.CliDriver.main(CliDriver.java:615)
```

```
        at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)
        at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:57)
        at
sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)
        at java.lang.reflect.Method.invoke(Method.java:606)
        at org.apache.hadoop.util.RunJar.run(RunJar.java:221)
        at org.apache.hadoop.util.RunJar.main(RunJar.java:136)
FAILED: ParseException line 1:7 cannot recognize input near 'create' 'databse' 'empdb' in ddl
statement
hive> create database empdb;
OK
Time taken: 1.882 seconds
hive> use empdb
    > create table emp ( emp_id int,name string,dept string,sal double join_date date );
FAILED: ParseException line 2:0 missing EOF at 'create' near 'empdb'
hive> create table emp ( emp_id int,name string,dept string,sal double, join_date date );
OK
Time taken: 0.543 seconds
hive> create external table dept ( dept_id int,dept_name string ) row format delimited fields
terminated by ',' stored as textfile location "/home/cloudera/Desktop/mock";
OK
Time taken: 0.08 seconds
hive> insert into table emp values (101,"chetan","HR",70000,'2023-01-10');
Query ID = cloudera_20250325232424_a7b77c1c-198c-4f87-bc14-2279d46fb349
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1742969110750_0001, Tracking URL =
http://quickstart.cloudera:8088/proxy/application_1742969110750_0001/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1742969110750_0001
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2025-03-25 23:24:40,415 Stage-1 map = 0%,  reduce = 0%
2025-03-25 23:24:48,430 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 1.42 sec
MapReduce Total cumulative CPU time: 1 seconds 420 msec
Ended Job = job_1742969110750_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to: hdfs://quickstart.cloudera:8020/user/hive/warehouse/emp/.hive-staging_hive_2025-
03-25_23-24-26_272_6369229161977442392-1/-ext-10000
Loading data to table default.emp
Table default.emp stats: [numFiles=1, numRows=1, totalSize=33, rawDataSize=32]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1   Cumulative CPU: 1.42 sec   HDFS Read: 4203 HDFS Write: 100
SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 420 msec
OK
Time taken: 24.736 seconds
hive>
    > insert into table emp values (102,"shailesh","HR",80000,'2024-01-10');
Query ID = cloudera_20250325233232_7c709046-641c-480f-a6e5-079c0b3d338a
Total jobs = 3
```

```
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1742969110750_0002, Tracking URL =
http://quickstart.cloudera:8088/proxy/application_1742969110750_0002/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1742969110750_0002
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2025-03-25 23:32:11,505 Stage-1 map = 0%,  reduce = 0%
2025-03-25 23:32:19,158 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 1.4 sec
MapReduce Total cumulative CPU time: 1 seconds 400 msec
Ended Job = job_1742969110750_0002
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to: hdfs://quickstart.cloudera:8020/user/hive/warehouse/emp/.hive-staging_hive_2025-
03-25_23-32-03_098_5454920136020912191-1/-ext-10000
Loading data to table default.emp
Table default.emp stats: [numFiles=2, numRows=2, totalSize=68, rawDataSize=66]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1   Cumulative CPU: 1.4 sec   HDFS Read: 4295 HDFS Write: 102 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 400 msec
OK
Time taken: 18.443 seconds
hive> insert into table emp values (103,"abhay","HR",100000,'2023-01-16');
Query ID = cloudera_20250325233232_2aae5917-4816-45de-a255-a0565b7d12f3
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1742969110750_0003, Tracking URL =
http://quickstart.cloudera:8088/proxy/application_1742969110750_0003/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1742969110750_0003
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2025-03-25 23:32:53,334 Stage-1 map = 0%,  reduce = 0%
2025-03-25 23:33:01,954 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 1.41 sec
MapReduce Total cumulative CPU time: 1 seconds 410 msec
Ended Job = job_1742969110750_0003
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to: hdfs://quickstart.cloudera:8020/user/hive/warehouse/emp/.hive-staging_hive_2025-
03-25_23-32-44_281_4286223111979945297-1/-ext-10000
Loading data to table default.emp
Table default.emp stats: [numFiles=3, numRows=3, totalSize=101, rawDataSize=98]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1   Cumulative CPU: 1.41 sec   HDFS Read: 4293 HDFS Write: 100
SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 410 msec
OK
Time taken: 18.992 seconds
hive> load data local inpath "/home/cloudera/Desktop/mock" into table emp;
FAILED: SemanticException Line 1:23 Invalid path '"/home/cloudera/Desktop/mock"': No files
matching path file:/home/cloudera/Desktop/mock
hive> load data local inpath "/home/cloudera/Desktop" into table emp;
```

FAILED: SemanticException Line 1:23 Invalid path "'/home/cloudera/Desktop'": source contains directory: file:/home/cloudera/Desktop/33343_elec_consumption
hive> load data local inpath "/home/cloudera/Desktop/mock.csv" into table emp;
FAILED: SemanticException Line 1:23 Invalid path "'/home/cloudera/Desktop/mock.csv'": No files matching path file:/home/cloudera/Desktop/mock.csv
hive> select e.emp_id,e.name from emp e join dept d on e.dept=d.dept_name;
Query ID = cloudera_20250325233838_50d88ad1-2f31-4547-a943-8076dbf4382c
Total jobs = 1
Execution log at: /tmp/cloudera/cloudera_20250325233838_50d88ad1-2f31-4547-a943-8076dbf4382c.log
2025-03-25 11:38:45 Starting to launch local task to process map join; maximum memory = 1013645312
2025-03-25 11:38:46 Dump the side-table for tag: 1 with group count: 0 into file: file:/tmp/cloudera/0a5e7521-348c-41bf-ae78-5f1d7806ae72/hive_2025-03-25_23-38-37_416_3673945639650006159-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile01--.hashtable
2025-03-25 11:38:46 Uploaded 1 File to: file:/tmp/cloudera/0a5e7521-348c-41bf-ae78-5f1d7806ae72/hive_2025-03-25_23-38-37_416_3673945639650006159-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile01--.hashtable (260 bytes)
2025-03-25 11:38:46 End of local task; Time Taken: 1.658 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1742969110750_0004, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1742969110750_0004/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1742969110750_0004
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2025-03-25 23:38:58,361 Stage-3 map = 0%, reduce = 0%
2025-03-25 23:39:08,180 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 1.22 sec
MapReduce Total cumulative CPU time: 1 seconds 220 msec
Ended Job = job_1742969110750_0004
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 1.22 sec HDFS Read: 6127 HDFS Write: 0 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 220 msec
OK
Time taken: 31.849 seconds
hive> insert into table dept values (1,"HR");
Query ID = cloudera_20250325234040_0d9e42d5-12d5-4bc6-86f0-239e31fba627
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1742969110750_0005, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1742969110750_0005/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1742969110750_0005
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2025-03-25 23:41:00,919 Stage-1 map = 0%, reduce = 0%
2025-03-25 23:41:10,881 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.46 sec
MapReduce Total cumulative CPU time: 1 seconds 460 msec
Ended Job = job_1742969110750_0005
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.

```
Moving data to: hdfs://quickstart.cloudera:8020/home/cloudera/Desktop/mock/.hive-staging_hive_2025-03-25_23-40-49_459_5434974939689600585-1/-ext-10000
Loading data to table default.dept
Table default.dept stats: [numFiles=1, numRows=1, totalSize=5, rawDataSize=4]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1   Cumulative CPU: 1.46 sec   HDFS Read: 3506 HDFS Write: 73 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 460 msec
OK
Time taken: 22.961 seconds
hive> insert into table dept values ( 2,"Engineering");
Query ID = cloudera_20250325234141_f19c4af8-a83c-4197-ba57-301bd5142a6e
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1742969110750_0006, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1742969110750_0006/
Kill Command = /usr/lib/hadoop/bin/hadoop job  -kill job_1742969110750_0006
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2025-03-25 23:41:35,784 Stage-1 map = 0%,  reduce = 0%
2025-03-25 23:41:45,540 Stage-1 map = 100%,  reduce = 0%, Cumulative CPU 1.44 sec
MapReduce Total cumulative CPU time: 1 seconds 440 msec
Ended Job = job_1742969110750_0006
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to: hdfs://quickstart.cloudera:8020/home/cloudera/Desktop/mock/.hive-staging_hive_2025-03-25_23-41-24_299_2302768609236652328-1/-ext-10000
Loading data to table default.dept
Table default.dept stats: [numFiles=2, numRows=2, totalSize=19, rawDataSize=17]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1   Cumulative CPU: 1.44 sec   HDFS Read: 3603 HDFS Write: 82 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 440 msec
OK
Time taken: 22.736 seconds
hive> select * from emp;
OK
101     chetan HR      70000.0         2023-01-10
102     shailesh        HR      80000.0         2024-01-10
103     abhay  HR      100000.0        2023-01-16
Time taken: 0.085 seconds, Fetched: 3 row(s)
hive> select * from dept;
OK
1       HR
2       Engineering
Time taken: 0.077 seconds, Fetched: 2 row(s)
hive> select e.emp_id,e.name from emp e join dept d on e.dept=d.dept_name;
Query ID = cloudera_20250325234242_dd5bbd75-5229-4b9f-b99c-d7f4f3363428
Total jobs = 1
Execution log at: /tmp/cloudera/cloudera_20250325234242_dd5bbd75-5229-4b9f-b99c-d7f4f3363428.log
2025-03-25 11:42:20  Starting to launch local task to process map join;     maximum memory = 1013645312
```

2025-03-25 11:42:22 Dump the side-table for tag: 1 with group count: 2 into file: file:/tmp/cloudera/0a5e7521-348c-41bf-ae78-5f1d7806ae72/hive_2025-03-25_23-42-13_744_6265123687685612945-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile11--.hashtable
2025-03-25 11:42:22 Uploaded 1 File to: file:/tmp/cloudera/0a5e7521-348c-41bf-ae78-5f1d7806ae72/hive_2025-03-25_23-42-13_744_6265123687685612945-1/-local-10003/HashTable-Stage-3/MapJoin-mapfile11--.hashtable (309 bytes)
2025-03-25 11:42:22 End of local task; Time Taken: 1.993 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1742969110750_0007, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1742969110750_0007/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1742969110750_0007
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2025-03-25 23:42:34,532 Stage-3 map = 0%, reduce = 0%
2025-03-25 23:42:44,048 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 1.44 sec
MapReduce Total cumulative CPU time: 1 seconds 440 msec
Ended Job = job_1742969110750_0007
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 1.44 sec HDFS Read: 6267 HDFS Write: 34 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 440 msec
OK
101 chetan
102 shailesh
103 abhay
Time taken: 31.38 seconds, Fetched: 3 row(s)
hive> select year(join_date) as year, max(sal) as max_sal from emp group by year(join_date);
Query ID = cloudera_20250325234545_190d60b5-e0b8-4513-b75c-de597fa38618
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1742969110750_0008, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1742969110750_0008/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1742969110750_0008
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2025-03-25 23:45:26,204 Stage-1 map = 0%, reduce = 0%
2025-03-25 23:45:36,019 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.29 sec
2025-03-25 23:45:47,806 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 2.62 sec
MapReduce Total cumulative CPU time: 2 seconds 620 msec
Ended Job = job_1742969110750_0008
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 2.62 sec HDFS Read: 7640 HDFS Write: 27 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 620 msec
OK

```
2023    100000.0
2024    80000.0
Time taken: 34.087 seconds, Fetched: 2 row(s)
Hbase:

                (root) at /usr/lib/hbase/bin/../bin/hirb.rb:210
[cloudera@quickstart ~]$ hbase shell
2025-03-26 00:30:44,216 INFO  [main] Configuration.deprecation: hadoop.native.lib is deprecated.
Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.0.0-cdh5.4.2, rUnknown, Tue May 19 17:07:29 PDT 2015

hbase(main):001:0> create_namespace 'EmployeeDB'

ERROR: org.apache.hadoop.hbase.NamespaceExistException: EmployeeDB
        at
org.apache.hadoop.hbase.master.TableNamespaceManager.create(TableNamespaceManager.java:15
0)
        at
org.apache.hadoop.hbase.master.TableNamespaceManager.create(TableNamespaceManager.java:12
5)
        at org.apache.hadoop.hbase.master.HMaster.createNamespace(HMaster.java:2083)
        at
org.apache.hadoop.hbase.master.MasterRpcServices.createNamespace(MasterRpcServices.java:393
)
        at
org.apache.hadoop.hbase.protobuf.generated.MasterProtos$MasterService$2.callBlockingMethod(
MasterProtos.java:44277)
        at org.apache.hadoop.hbase.ipc.RpcServer.call(RpcServer.java:2035)
        at org.apache.hadoop.hbase.ipc.CallRunner.run(CallRunner.java:107)
        at org.apache.hadoop.hbase.ipc.RpcExecutor.consumerLoop(RpcExecutor.java:130)
        at org.apache.hadoop.hbase.ipc.RpcExecutor$1.run(RpcExecutor.java:107)
        at java.lang.Thread.run(Thread.java:745)

Here is some help for this command:
Create namespace; pass namespace name,
and optionally a dictionary of namespace configuration.
Examples:

  hbase> create_namespace 'ns1'
  hbase> create_namespace 'ns1', {'PROERTY_NAME'=>'PROPERTY_VALUE'}


hbase(main):002:0> create 'EmployeeDB:Department', {NAME => 'dept_details', VERSIONS =>
1}

0 row(s) in 1.4280 seconds

=> Hbase::Table - EmployeeDB:Department
hbase(main):003:0>
hbase(main):004:0* put 'EmployeeDB:Employee', '103', 'emp_details:name', 'Chetan';
```

```
hbase(main):005:0* PUT 'EmployeeDB:Employee', '103', 'emp_details:department', 'Finance';
hbase(main):006:0* PUT 'EmployeeDB:Employee', '103', 'emp_details:salary', '85000';
hbase(main):007:0* PUT 'EmployeeDB:Employee', '103', 'emp_details:joining_date', '2021-08-
22';PUT 'EmployeeDB:Employee', '103', 'emp_details:joining_date', '2021-08-22';
[2]+  Stopped                 hbase shell
[cloudera@quickstart ~]$ hbase shell
2025-03-26 00:33:14,095 INFO  [main] Configuration.deprecation: hadoop.native.lib is deprecated.
Instead, use io.native.lib.available
HBase Shell; enter 'help<RETURN>' for list of supported commands.
Type "exit<RETURN>" to leave the HBase Shell
Version 1.0.0-cdh5.4.2, rUnknown, Tue May 19 17:07:29 PDT 2015

hbase(main):001:0> create 'EmployeeDB:Department', {NAME => 'dept_details', VERSIONS =>
1}

ERROR: Table already exists: EmployeeDB:Department!

Here is some help for this command:
Creates a table. Pass a table name, and a set of column family
specifications (at least one), and, optionally, table configuration.
Column specification can be a simple string (name), or a dictionary
(dictionaries are described below in main help output), necessarily
including NAME attribute.
Examples:

Create a table with namespace=ns1 and table qualifier=t1
  hbase> create 'ns1:t1', {NAME => 'f1', VERSIONS => 5}

Create a table with namespace=default and table qualifier=t1
  hbase> create 't1', {NAME => 'f1'}, {NAME => 'f2'}, {NAME => 'f3'}
  hbase> # The above in shorthand would be the following:
  hbase> create 't1', 'f1', 'f2', 'f3'
  hbase> create 't1', {NAME => 'f1', VERSIONS => 1, TTL => 2592000, BLOCKCACHE => true}
  hbase> create 't1', {NAME => 'f1', CONFIGURATION => {'hbase.hstore.blockingStoreFiles' =>
'10'}}

Table configuration options can be put at the end.
Examples:

  hbase> create 'ns1:t1', 'f1', SPLITS => ['10', '20', '30', '40']
  hbase> create 't1', 'f1', SPLITS => ['10', '20', '30', '40']
  hbase> create 't1', 'f1', SPLITS_FILE => 'splits.txt', OWNER => 'johndoe'
  hbase> create 't1', {NAME => 'f1', VERSIONS => 5}, METADATA => { 'mykey' => 'myvalue' }
  hbase> # Optionally pre-split the table into NUMREGIONS, using
  hbase> # SPLITALGO ("HexStringSplit", "UniformSplit" or classname)
  hbase> create 't1', 'f1', {NUMREGIONS => 15, SPLITALGO => 'HexStringSplit'}
  hbase> create 't1', 'f1', {NUMREGIONS => 15, SPLITALGO => 'HexStringSplit',
REGION_REPLICATION => 2, CONFIGURATION =>
{'hbase.hregion.scan.loadColumnFamiliesOnDemand' => 'true'}}

You can also keep around a reference to the created table:
```

```
hbase> t1 = create 't1', 'f1'
```

Which gives you a reference to the table named 't1', on which you can then call methods.

```
hbase(main):002:0>
```